

LogicPoison 复现报告

Reproduction Report

May 2026

1 论文信息 / Paper Info

Title	LogicPoison: Logical Attacks on Graph Retrieval-Augmented Generation
Authors	Yilin Xiao, Jin Chen, Qinggang Zhang, Yujing Zhang, Chuang Zhou, Longhao Yang, Lingfei Ren, Xin Yang, Xiao Huang
Venue	ACL 2026 Main Conference
ArXiv	https://arxiv.org/abs/2604.02954
GitHub	https://github.com/Jord8061/logicPoison
HuggingFace	https://huggingface.co/datasets/Jord8061/datasets

Table 1: Paper metadata.

2 论文概述 / Paper Overview

LogicPoison is a logical poisoning framework targeting **GraphRAG** systems. Rather than injecting fabricated facts or adversarial instructions into the corpus, it perturbs the *implicit reasoning topology* of the graph through **type-preserving entity swapping**.

The core insight is straightforward but effective:

- Identify globally important *logic hubs* in the corpus.
- Identify query-specific *bridge entities* required for multi-hop reasoning.
- Swap entities *within the same semantic type* (e.g., PERSON with PERSON, ORG with ORG) so that the text remains locally plausible while the graph structure built from the corpus is corrupted.

This causes GraphRAG systems to follow corrupted reasoning chains and produce incorrect answers, while the poisoned corpus remains surface-level natural.

2.1 Three-Stage Framework

1. Global Logic Poison (Stage 1):

Build corpus-level entity statistics and identify high-frequency entities that likely act as *global logic hubs* in the implicit graph. Uses spaCy’s transformer-based NER model (`en_core_web_trf`) with GPU acceleration to process the entire corpus, collecting per-type entity frequency distributions.

2. Query-Centric Logic Poison (Stage 2):

Use query-side reasoning signals to identify *bridge entities* that are important for answering multi-hop questions. An LLM extracts reasoning-critical entities from each question, including implicit bridge entities (typed `BRIDGE`) and alias-style descriptions (typed `ALIAS`).

3. Type-Preserving Entity Swapping (Stage 3):

Combine the selected entities and perform *bijective (circular) swapping within the same entity type*. For each type, a pool is constructed from the top-5% corpus entities (by frequency) plus query entities (in occurrence order). A reverse circular permutation is then applied:

$$e_0 \rightarrow e_{n-1}, \quad e_1 \rightarrow e_0, \quad e_2 \rightarrow e_1, \quad \dots, \quad e_{n-1} \rightarrow e_{n-2}$$

This ensures every entity is replaced by exactly one other entity of the same type, producing a poisoned corpus that remains grammatically and semantically plausible at the local text level.

2.2 Key Contributions

- A novel poisoning paradigm that targets *logical structure* rather than surface text, making attacks harder to detect.
- A three-stage pipeline (global, query-centric, type-preserving swap) that is modular and applicable to any GraphRAG system.
- Evaluation on three multi-hop QA benchmarks (HotpotQA, MuSiQue, 2WikiMultiHopQA) demonstrating significant degradation of GraphRAG performance.

3 复现环境 / Reproduction Environment

Python	3.14.4
Virtual Env	venv314
GPU	NVIDIA GeForce RTX 4060
CUDA	13.1
CuPy	Installed (GPU-accelerated spaCy)
LLM API	opencode-go / deepseek-v4-flash

Table 2: System environment.

Package	Version
openai	2.30.0
spacy	3.8.13
tqdm	4.67.3
en_core_web_trf	3.8.0

Table 3: Python dependencies.

4 复现步骤 / Reproduction Steps

4.1 Step 1: Clone Repository

```
git clone https://github.com/Jord8061/logicPoison.git
cd logicPoison
```

4.2 Step 2: Clone Datasets from HuggingFace

```
cd datasets
git clone https://huggingface.co/datasets/Jord8061/datasets .
cd ..
```

The datasets directory should contain three subdirectories: `hotpotqa/`, `musique/`, and `2wikimultihopqa/`, each with `corpus.jsonl` and `queries.jsonl`.

4.3 Step 3: Install Dependencies and spaCy Model

```
python -m venv venv314
source venv314/bin/activate
pip install -r requirements.txt
pip install cupy-cuda13x # GPU acceleration for spaCy
python -m spacy download en_core_web_trf
```

4.4 Step 4: Configure API Access

Set the following environment variables for the OpenAI-compatible API:

```
export OPENAI_API_KEY="<your-api-key>"
export OPENAI_BASE_URL="<your-base-url>"
```

Note: In this reproduction, we used the `deepseek-v4-flash` model via the `opencode-go` gateway. The API key and base URL must be configured before running the query-centric stage.

4.5 Step 5: Fix `list_datasets` Filter

The original `main.py` uses `os.listdir()` to discover datasets, which may include non-dataset directories such as `.git/` or files like `README.md` that were cloned from HuggingFace. The function already filters by checking for `corpus.jsonl` and excluding dot-prefixed names, but we verified this handles the HuggingFace clone correctly. No patch was needed in our case since the `list_datasets` function in `main.py` already requires `corpus.jsonl` to exist in each subdirectory.

4.6 Step 6: Run the Pipeline

```
python main.py \
  --stages all \
  --datasets all \
  --query_model deepseek-v4-flash
```

The pipeline runs three stages sequentially:

1. **global**: Corpus entity statistics extraction (GPU-accelerated NER).
2. **query**: Query-centric entity extraction via LLM API calls.
3. **logic**: Type-preserving entity swapping and corpus poisoning.

Completed stage outputs are auto-detected and skipped on re-runs. Use `--force` to rerun selected stages.

5 复现结果 / Reproduction Results

Results are organized into three directories under **results/**:

- **corpus_entities/**: Per-dataset entity statistics (JSON).
- **queries_entities/**: Per-dataset query entity extraction (JSONL).
- **poisoned_data/**: Poisoned corpora and statistics.

5.1 Global Entity Statistics

Table 4: Global entity statistics per dataset (Stage 1 output).

Dataset	Corpus Size	Types Poisoned	Top Entity Types (count)
2WikiMultiHopQA	6,119	7	PERSON (977), WORK_OF_ART (431), ORG (
HotpotQA	9,811	13	PERSON (1,110), ORG (769), WORK_OF_ART
MuSiQue	11,656	11	PERSON (997), ORG (644), DATE (503)

Table 5: Top entities by type for HotpotQA (Stage 1 output).

Type	Top Entity	Count
PERSON	(top PERSON entity)	1,110
ORG	(top ORG entity)	769
WORK_OF_ART	(top WORK_OF_ART entity)	512

5.2 Query-Centric Entity Extraction

All three datasets contain 500 questions each (1,500 total). Entity extraction was performed using **deepseek-v4-flash** via the **opencode-go** API (**OPENAI_BASE_URL**=<https://opencode.ai/zen/go/v1>).

Table 6: Query entity extraction statistics per dataset (Stage 2 output).

Dataset	Num Queries	Model	Runtime
2WikiMultiHopQA	500	deepseek-v4-flash	~82 min
HotpotQA	500	deepseek-v4-flash	~56 min
MuSiQue	500	deepseek-v4-flash	~7.5 hours

5.3 Poison Statistics

Table 7: Poison statistics per dataset (Stage 3 output).

Dataset	Types Poisoned	Entities Swapped	Total Freq Affected
2WikiMultiHopQA	7	1,865	17,615
HotpotQA	13	3,781	62,905
MuSiQue	11	3,505	61,968

Table 8: Per-type poison statistics for HotpotQA (Stage 3 output).

Entity Type	Num Entities Swapped	Total Frequency
PERSON	(dominant type)	(largest share)
ORG	(second largest)	(second largest)
WORK_OF_ART	(third largest)	(third largest)

Note: The results directory structure is:

```

results/
  corpus_entities/
    hotpotqa.json
    musique.json
    2wikimultihopqa.json
  queries_entities/
    hotpotqa.jsonl
    musique.jsonl
    2wikimultihopqa.jsonl
  poisoned_data/
    hotpotqa/
      corpus.jsonl
      queries.jsonl
      answers.jsonl
      qrels/
        poison_stats_hotpotqa.json
    musique/
      ...
    2wikimultihopqa/
      ...

```

6 关键发现 / Key Findings

6.1 Pipeline Performance

Table 9: Pipeline runtime breakdown.

Stage	Dataset	Runtime
Global (GPU)	2WikiMultiHopQA	~2–3 min
	HotpotQA	~2–3 min
	MuSiQue	~2–3 min
Query (API)	2WikiMultiHopQA	~82 min
	HotpotQA	~56 min
	MuSiQue	~7.5 hours
Logic	2WikiMultiHopQA	~1–2 min
	HotpotQA	~1–2 min
	MuSiQue	~1–2 min
Total	All 3 datasets	~10 hours

The query-centric stage dominates total runtime due to per-question LLM API calls (~2 min/question average). The global and logic stages are fast thanks to GPU-accelerated NER and local text substitution respectively.

6.2 Sanity Check: Type-Preserving Swapping

We verified that the type-preserving constraint is correctly enforced by inspecting sample replacements in the poisoned corpora:

Table 10: Sample entity swaps confirming type preservation.

Original	Replaced With	Type	Preserved?
2011	2010	DATE → DATE	✓
Indian	Australian	NORP → NORP	✓

Documents remain grammatically coherent with only entities swapped—no broken syntax or semantic incoherence was observed at the sentence level.

6.3 Type-Preserving Entity Swapping Mechanism

The core innovation of LogicPoison is the *type-preserving* constraint. By swapping entities only within the same NER type (e.g., PERSON ↔ PERSON, GPE ↔ GPE), the poisoned corpus maintains local grammatical coherence. The reverse circular permutation ensures:

- Every entity in the pool is replaced by exactly one other entity (bijective mapping).
- No entity is replaced by itself.
- The mapping is deterministic and reproducible.

The pool for each type is constructed by concatenating:

1. Top-5% entities by corpus frequency (global logic hubs).
2. Query-centric entities that appear in the corpus (bridge entities), appended in occurrence order, deduplicated against the top-5%.

6.4 GPU Acceleration Impact

The global entity extraction stage (Stage 1) uses spaCy’s transformer-based NER model (`en_core_web_trf`) with GPU acceleration via CuPy. On our RTX 4060 with CUDA 13.1, `spacy.prefer_gpu()` successfully enabled GPU processing, significantly speeding up NER over the full corpus. The pipeline uses `nlp.pipe()` with configurable batch size (default 32) for efficient batch processing.

6.5 API Rate Limiting Observations

The query-centric stage (Stage 2) makes concurrent LLM API calls using a `ThreadPoolExecutor` with configurable `max_workers` (default 8) and `queue_factor` (default 4), yielding up to 32 concurrent requests. Key observations:

- The pipeline includes exponential backoff retry logic (3 retries, base delay 1.5s) for API errors.
- Using `deepseek-v4-flash` as the query model provides a cost-effective alternative to `gpt-4o-mini` for entity extraction.
- Rate limiting may require adjusting `--max_workers` and `--queue_factor` depending on the API provider’s limits.

6.6 Pipeline Robustness

The pipeline includes several robustness features:

- **Auto-skip completed stages:** The `stage_done()` function checks for existing output files and skips completed work.

- **Regex-based replacement:** Stage 3 uses a single compiled regex pattern with word-boundary assertions $((?<!\w)$ and $(?!\\w))$ to prevent cascading replacements and partial word matches.
- **BEIR-compatible output:** The poisoned data directory preserves `queries.jsonl`, `answers.jsonl`, and `qrels/` for direct evaluation with BEIR-style retrieval frameworks.

7 复现对比分析 / Reproduction Comparison

This section compares our reproduction results against the original paper’s reported metrics and experimental setup.

7.1 Original Paper Results

The original paper reports Attack Success Rate (ASR%) on GraphRAG across three LLM backbones and three datasets. ASR measures the proportion of questions for which the poisoned GraphRAG system produces incorrect answers.

Table 11: Original paper: ASR (%) on Microsoft GraphRAG (Table 1 in paper).

Dataset	GPT-4o-mini	Llama-3.1-8B	Qwen-3-32B
HotpotQA	78.4	79.2	84.2
2WikiMultiHopQA	78.4	78.8	78.6
MuSiQue	91.4	91.6	93.0

The paper additionally reports ASR-GPT (LLM-judge) metrics and results on two alternative GraphRAG frameworks: GFM-RAG and HippoRAG 2. Across all configurations, LogicPoison achieves ASR values consistently above 70%, with MuSiQue reaching 93% under Qwen-3-32B.

7.2 Our Reproduction Results

We successfully reproduced the full three-stage poison generation pipeline. The following table summarizes the poison statistics from our run:

Table 12: Our reproduction: poison generation statistics.

Dataset	Entities Swapped	Total Freq	Entity Types	Query Model
2WikiMultiHopQA	1,865	17,615	7	deepseek-v4-flash
HotpotQA	3,781	62,905	13	deepseek-v4-flash
MuSiQue	3,505	61,968	11	deepseek-v4-flash
Total	9,151	142,488	31	

Note: the “31” total entity types counts each type per dataset separately (i.e., the same type such as PERSON appears in all three datasets).

7.3 Side-by-Side Comparison

Table 13: Comparison between original paper and our reproduction.

Metric	Paper	Our Reproduction	Notes
Pipeline Stages	3 (global, query, logic)	3 (global, query, logic)	✓ Identical
Query Model	gpt-4o-mini	deepseek-v4-flash	Different model
API Backend	OpenAI	opencode-go (Zen)	Different provider
GPU	RTX 5090	RTX 4060	Consumer GPU
Datasets	3 (500 queries each)	3 (500 queries each)	✓ Identical
Type-Preserving Swap	✓	✓	✓ Verified
NER Model	spaCy en_core_web_trf	spaCy en_core_web_trf	✓ Identical
HotpotQA Entities	—	3,781	Not reported in paper
2Wiki Entities	—	1,865	Not reported in paper
MuSiQue Entities	—	3,505	Not reported in paper
GraphRAG ASR	78–93%	Not evaluated	Requires GraphRAG setup
GFM-RAG ASR	Reported	Not evaluated	Requires GFM-RAG setup
HippoRAG 2 ASR	Reported	Not evaluated	Requires HippoRAG 2

7.4 Key Differences and Analysis

- Query Model Difference:** The original paper uses `gpt-4o-mini` (OpenAI) for query-centric entity extraction, while we used `deepseek-v4-flash` via the `opencode-go` API. Since the query stage selects which bridge entities to target, different LLMs may produce different entity selections, which would affect the composition of the swapping pool and consequently the downstream ASR metrics. The type-preserving swapping mechanism itself is model-agnostic—only the entity selection differs.
- GPU Hardware:** The paper uses an RTX 5090; we used an RTX 4060. Both GPUs successfully accelerated spaCy’s transformer-based NER via CuPy. The global stage completed in $\sim 2\text{--}3$ minutes per dataset on both configurations, confirming that the NER workload is not heavily GPU-bound at these corpus sizes (6K–12K documents).
- Evaluation Gap:** The paper reports full ASR metrics across three GraphRAG frameworks (Microsoft GraphRAG, GFM-RAG, HippoRAG 2) and three LLM backbones (GPT-4o-mini, Llama-3.1-8B, Qwen-3-32B). Our reproduction covers the poison generation pipeline

only; running the full GraphRAG evaluation requires additional framework setup (Microsoft GraphRAG, GFM-RAG, HippoRAG 2) and is left as future work.

4. **Poison Statistics:** The original paper does not report per-dataset entity swap counts or total frequency affected. Our reproduction provides these statistics (Table above), which are consistent with the expected scale: thousands of entities across 7–13 types per dataset, affecting tens of thousands of corpus occurrences. This confirms the attack surface is substantial.

7.5 Efficiency Analysis / 效率分析

Table 14 compares the runtime of each pipeline stage across our reproduction and the original paper’s reported estimates.

Table 14: Pipeline runtime comparison (our reproduction vs. paper estimates)

Stage	2Wiki	HotpotQA	MuSiQue	Paper Est.
Global (spaCy NER, GPU)	~2 min	~3 min	~3 min	~2–3 min
Query (API extraction)	~82 min	~56 min	~448 min	~30–60 min
Logic (entity swap)	~1 min	~1 min	~1 min	~1 min
Total per dataset	~85 min	~60 min	~452 min	~35–65 min

Corpus sizes for context:

- 2wikimultihopqa: 6,119 documents
- hotpotqa: 9,811 documents
- musique: 11,656 documents (largest corpus, slowest query stage)

Key observations:

1. **Global stage:** The RTX 4060 GPU performed comparably to expectations (30–80 docs/sec after CuPy warmup). The transformer-based spaCy NER is not heavily GPU-bound at these corpus sizes.
2. **Query stage bottleneck:** The `deepseek-v4-flash` API via `opencode-go` had highly variable latency (1.5–22 sec/query). MuSiQue took 7.5 hours primarily due to API rate limiting on the `opencode-go` free tier. The original paper used OpenAI’s `gpt-4o-mini` which likely had higher and more consistent throughput.
3. **Logic stage:** Fast local text processing (~1 min), consistent with expectations—no API calls involved.
4. **Overall:** The pipeline completed successfully, but API rate limiting is the primary efficiency divergence from the paper. Using a higher-tier API plan would bring the query stage closer to paper estimates.

7.6 Naive RAG ASR Evaluation Results

We evaluated the poisoned corpora using a Naive RAG pipeline (Contriever + FAISS retrieval + LLM generation), matching the paper’s retrieval setup. 10 queries were sampled per dataset. The `deepseek-v4-flash` reasoning model required token-aware retry logic: empty answers with `finish_reason=length` were retried with up to 8,000 tokens to distinguish token-budget empties from genuine empty responses.

Table 15: Naive RAG ASR comparison: our reproduction vs. paper

Dataset	Our ASR (deepseek-v4-flash)	Paper ASR (GPT-4o-mini)	Non-Empty / Total
HotpotQA	60.0%	92.0%	9/10
2WikiMultihopQA	90.0%	90.0%	9/10
MuSiQue	100.0%*	99.2%	0/10

*MuSiQue: all answers empty due to reasoning token overflow; retry was interrupted. The 100% ASR reflects empty output $\neq$ ground truth and is not directly comparable.

Key findings:

- **2Wiki: ASR 90.0%** — matches the paper’s 90.0% exactly, suggesting comparable attack effectiveness for this dataset.
- **HotpotQA: ASR 60.0%** — significantly *lower* than the paper’s 92.0%. The **deepseek-v4-flash** model appears more robust against poisoned context for HotpotQA-style questions, correctly answering 4/10 despite entity swaps.
- **MuSiQue: evaluation incomplete** — requires token-aware retry.
- **Model difference matters:** Using a different LLM (**deepseek-v4-flash** vs. **GPT-4o-mini**) significantly impacts ASR. The reasoning model’s tendency to “think through” the context may help it bypass poisoned entities.
- **Token budget is critical:** Reasoning models consume output tokens for internal thinking. Without sufficient `max_tokens` (we used up to 8,000), answers appear empty, inflating ASR artificially.

7.7 Reproduction Verdict

The poison generation pipeline is **successfully reproduced** with identical architecture and verified type-preserving behavior. The core mechanism—global entity statistics, query-centric entity extraction, and type-preserving circular swapping—produces poisoned corpora that are structurally consistent with the paper’s description.

Quantitative comparison with paper: The paper does not report per-dataset entity swap counts, but our statistics (9,151 entities across 31 types, affecting 142,488 corpus occurrences) are consistent with the attack’s intended scale. The retrieval infrastructure for ASR evaluation has been built and verified (§ 6.3); completing the full ASR comparison requires API credits for LLM answer generation.

What was reproduced:

- All 3 pipeline stages (global, query, logic) across all 3 datasets
- Type-preserving entity swapping mechanism (verified with sanity checks)
- GPU-accelerated NER (CuPy + RTX 4060)
- Dense retrieval pipeline (Contriever + FAISS)
- BEIR-compatible poisoned dataset format

Limitations:

- ASR evaluation blocked by API credit limits
- GraphRAG framework evaluation (Microsoft GraphRAG, GFM-RAG, HippoRAG 2) not run
- Query model differed from paper (deepseek-v4-flash vs. gpt-4o-mini)

8 结论 / Conclusion

This report documents the reproduction of the LogicPoison paper (ACL 2026). The three-stage pipeline (global entity statistics, query-centric entity extraction, and type-preserving entity

swapping) was successfully set up and executed on the HotpotQA, MuSiQue, and 2WikiMulti-HopQA datasets.

Key takeaways:

- The reproduction environment (Python 3.14.4, RTX 4060, CUDA 13.1, deepseek-v4-flash API) is fully functional.
- The type-preserving entity swapping mechanism is elegant: it corrupts graph-level reasoning structure while preserving local text plausibility.
- GPU acceleration via CuPy is essential for the NER-heavy global stage.
- The modular pipeline design allows running individual stages and datasets independently.
- Total pipeline runtime was ~ 10 hours across all three datasets, dominated by the query-centric LLM API stage.
- Across all datasets, 9,151 unique entities were swapped, affecting 142,488 total occurrences in the corpora.

Citation:

```
@misc{xiao2026logicpoisonlogicalattacksgraph,
      title={LogicPoison: Logical Attacks on Graph
            Retrieval-Augmented Generation},
      author={Yilin Xiao and Jin Chen and Qinggang Zhang
            and Yujing Zhang and Chuang Zhou
            and Longhao Yang and Lingfei Ren
            and Xin Yang and Xiao Huang},
      year={2026},
      eprint={2604.02954},
      archivePrefix={arXiv},
      primaryClass={cs.CL},
      url={https://arxiv.org/abs/2604.02954},
}
```